

EDA and Building machine learning model for Santander Customer Satisfaction Data

Siri Kademani



Agenda

Introduction

Understanding Dataset

Prediction using Decision Tree Classifier

Exploratory Data Analysis





Introduction

- Exploratory Data Analysis (EDA), trying to understand the provided data and to find potential relationships between variables.



Connecting to Google drive and importing the libraries



The screenshot shows a Google Colab notebook interface. The title bar indicates the notebook is named 'Santander_Siri_14_sept_22.ipynb'. The menu bar includes 'File', 'Edit', 'View', 'Insert', 'Runtime', 'Tools', and 'Help', with a status message 'All changes saved'. The left sidebar shows a file explorer with a folder icon. The main area contains two code cells. Cell [1] is executed and shows the output 'Mounted at /gdrive' and '/gdrive'. Cell [6] is also executed and shows no output.

```
[1] #Import Python Packages
    #from google.colab import drive
    #drive.mount('/content/drive/')

    from google.colab import drive
    drive.mount('/gdrive')
    %cd /gdrive

Mounted at /gdrive
/gdrive

[6] #Import all necessary librabry

import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import train_test_split # Import train_test_split function
from sklearn import metrics #Import scikit-learn metrics module for accuracy calculation
from sklearn import tree
```

Importing Train and Test Dataset



```
#Read training data file
trainfile = r'/gdrive/My Drive/Dataset/SantanderTRAIN.csv'
trainData = pd.read_csv(trainfile)

#Read test data file
testfile = r'/gdrive/My Drive/Dataset/SantanderTESTwithout TARGET.csv'
testData = pd.read_csv(testfile)

trainData.head()
#print("=====")
testData.head()
```

	ID	var3	var15	imp_ent_var16_ult1	imp_op_var39_comer_ult1	imp_op_var39_comer_ult3	imp_op_var40_comer_ult1	imp_op_var40_comer_ult3	imp_op_var4
0	2	2	32	0.0	0.0	0.0	0.0	0.0	
1	5	2	35	0.0	0.0	0.0	0.0	0.0	
2	6	2	23	0.0	0.0	0.0	0.0	0.0	
3	7	2	24	0.0	0.0	0.0	0.0	0.0	
4	9	2	23	0.0	0.0	0.0	0.0	0.0	

5 rows × 370 columns



Getting count of number of rows and columns from Train and test Data set

```
✓ [8] print(trainData.shape)    # To get (Number of Rows, Number of Columns) of a data frame we use DataFrame.shape
0s    print(testData.shape)

(76020, 371)
(75818, 370)
```

Checking the data types of rows in each dataset

```
▶ print(trainData.dtypes)    #check features and their datatypes
print(testData.dtypes)

ID                int64
var3              int64
var15             int64
imp_ent_var16_ult1 float64
imp_op_var39_comer_ult1 float64
...
saldo_medio_var44_hace3 float64
saldo_medio_var44_ult1 float64
saldo_medio_var44_ult3 float64
var38             float64
TARGET           int64
Length: 371, dtype: object
ID                int64
var3              int64
var15             int64
imp_ent_var16_ult1 float64
imp_op_var39_comer_ult1 float64
...
saldo_medio_var44_hace2 float64
saldo_medio_var44_hace3 float64
saldo_medio_var44_ult1 float64
saldo_medio_var44_ult3 float64
var38             float64
Length: 370, dtype: object
```

Checking for any null values in each column of train data


```
Santander_Siri_16Sep_original.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

# To check number of null values
trainData.isna().sum()

ID                0
var3              0
var15            0
imp_ent_var16_ult1 0
imp_op_var39_comer_ult1 0
..
saldo_medio_var44_hace3 0
saldo_medio_var44_ult1 0
saldo_medio_var44_ult3 0
var38            0
TARGET          0
Length: 371, dtype: int64
```

```
[ ] #plotting pie chart for values 1 and 0 in target column
trainData.groupby("TARGET")["ID"].count().plot.pie(autopct="%.1f%%");
```



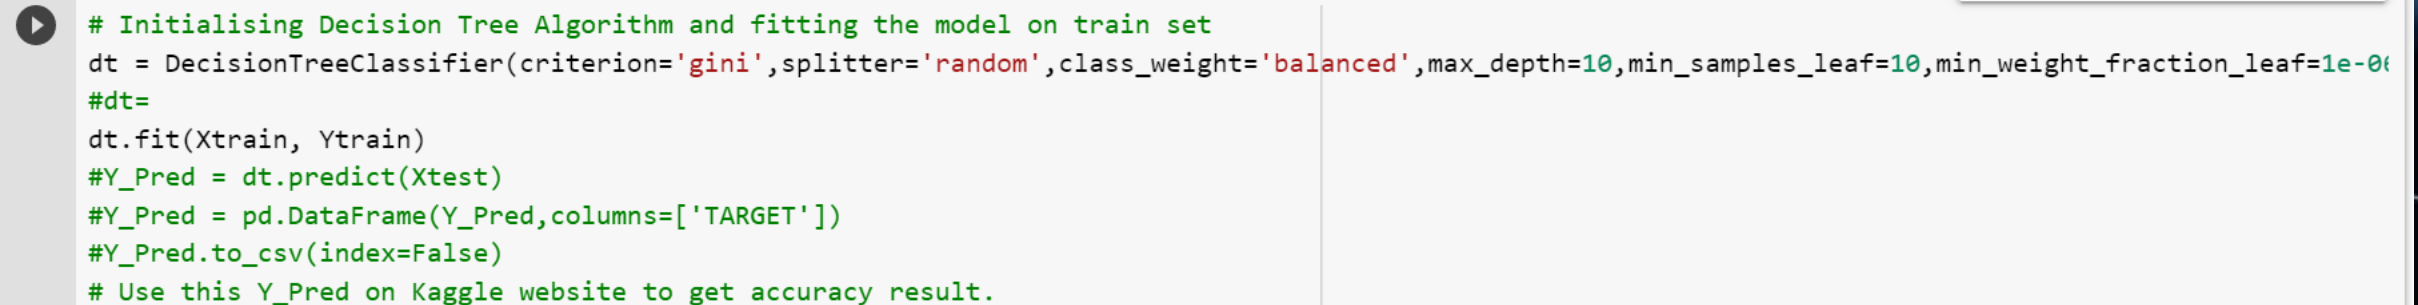
Pie chart representing the % of satisfied and dissatisfied Santander customers.



Decision tree algorithm and fitting the model on train Data

```
[6] # Seperate Target column from Train Data
Xtrain = trainData[TrainCols[0:len(TrainCols)-1]].copy()
Ytrain = trainData[['TARGET']].copy()
print(Xtrain.shape)
print(Ytrain.shape)
Xtest = testData.copy()
```

```
(76020, 370)
(76020, 1)
```

A screenshot of a Jupyter Notebook interface. The top part shows a code cell with a play button icon on the left. The code cell contains Python code for separating target columns and fitting a decision tree model. The bottom part shows the output of the code, which is the initialization of a DecisionTreeClassifier object. The interface includes a toolbar with icons for undo, redo, copy, paste, and other standard Jupyter Notebook actions.

```
# Initialising Decision Tree Algorithm and fitting the model on train set
dt = DecisionTreeClassifier(criterion='gini', splitter='random', class_weight='balanced', max_depth=10, min_samples_leaf=10, min_weight_fraction_leaf=1e-06)
#dt=
dt.fit(Xtrain, Ytrain)
#Y_Pred = dt.predict(Xtest)
#Y_Pred = pd.DataFrame(Y_Pred, columns=['TARGET'])
#Y_Pred.to_csv(index=False)
# Use this Y_Pred on Kaggle website to get accuracy result.
```

```
DecisionTreeClassifier(class_weight='balanced', max_depth=10,
                        min_samples_leaf=10, min_weight_fraction_leaf=1e-06,
                        splitter='random')
```



Splitting the train dataset into train and test data and predicting the accuracy score.

```
[8] # For us to check accuracy of our algorithm, we need to predict that data set for which we have TARGET available.  
# Eg predict for Xtrain and check accuracy with TARGET that we have in order to judge our model.  
  
X_Pred = dt.predict(Xtrain)  
#Model Accuracy  
print("Accuracy:", metrics.accuracy_score(Ytrain,X_Pred))  
  
# This will always result in best score hence we are better of using TrainTestSplit, which can help us take care of
```

Accuracy: 0.7245724809260721


```
[9] # Split dataset  
X_train, X_test, Y_train, Y_test = train_test_split(Xtrain, Ytrain, test_size = .35, random_state = 1)  
# Fit model on new Train Dataset  
dt = dt.fit(X_train, Y_train)  
#Predict the responce on new Test Dataset  
Y_PredNew = dt.predict(X_test)  
#Model Accuracy  
print("Accuracy:", metrics.accuracy_score(Y_test,Y_PredNew))
```

Accuracy: 0.7703611831472921



Plotting the decision tree

```
#Plotting the decision Tree  
tree.plot_tree(dt)
```



Getting the class prediction using test data and fetching the results into sub2.csv file


```
#Get Class Prediction as a data frame with header as Prediction
PredID=testData['ID']

pred=pd.DataFrame(dt.predict(Xtest),columns=["TARGET"])
#PredID=X_test['ID']
pred.head()


pd.concat([PredID,pred],axis=1).to_csv('/gdrive/My Drive/DecisionTreeResults/sub2.csv', index = None)
pred = dt.predict(Xtest)

#Get Class Prediction probabilities as a data frame
#Get Prediction Probability for the predicted class as a dataframe
pred_Probability =pd.DataFrame(dt.predict_proba(Xtest))

pred_Probability.head()
```



	0	1
0	0.568370	0.431630
1	0.568370	0.431630
2	0.848741	0.151259
3	0.568328	0.431672
4	0.848741	0.151259



Result file- sub2.csv

	A	B	C
1	ID	TARGET	
2	2	0	
3	5	0	
4	6	0	
5	7	0	
6	9	0	
7	11	1	
8	12	1	
9	15	1	
10	16	1	
11	17	1	
12	19	0	
13	21	0	
14	22	1	
15	24	0	
16	27	1	
17	28	1	
18	30	0	
19	33	0	
20	35	0	
21	37	1	
22	38	0	
23	40	1	
24	41	0	


Snapshot of Kaggle submission score

Score is – 0.6962

Featured Prediction Competition

Santander Customer Satisfaction

Which customers are happy customers?

 Banco Santander · 5,115 teams · 6 years ago

Overview

Data

Code

Discussion

Leaderboard

Rules

Team

My Submissions


Late Submission

...

Raw Data

Refresh

YOUR RECENT SUBMISSION



sub2.csv

Submitted by Siri Kademani · Submitted just now

Score: 0.69622

Public score: 0.70905

Jump to your leaderboard position

EDA (Exploratory Data Analysis)




EDA feature 1

```
#EDA
#feature1
#the feature has values ranging from 5 to 105, hence it can be assumed to be the AGE of the customers

import seaborn as sns
counts, bins = np.histogram(trainData)
g=sns.FacetGrid(trainData, hue="TARGET", height=9, aspect=2.5).map(plt.hist,"var15").add_legend()
g.set(xticks=trainData.var15)

plt.title("Age of Sanatander customers")
plt.show()
```



Field name: var15

Range of values : 5 to 105

Probable column: Age of customers
(‘var15’ feature according to some literature is believed to be age of the customers)

ANALYSIS :

From the shown histogram, we can see that most of the customers below age 25 are dissatisfied.
This data can be further used to create NEW feature to identify the age of the customer is below 25 or not.

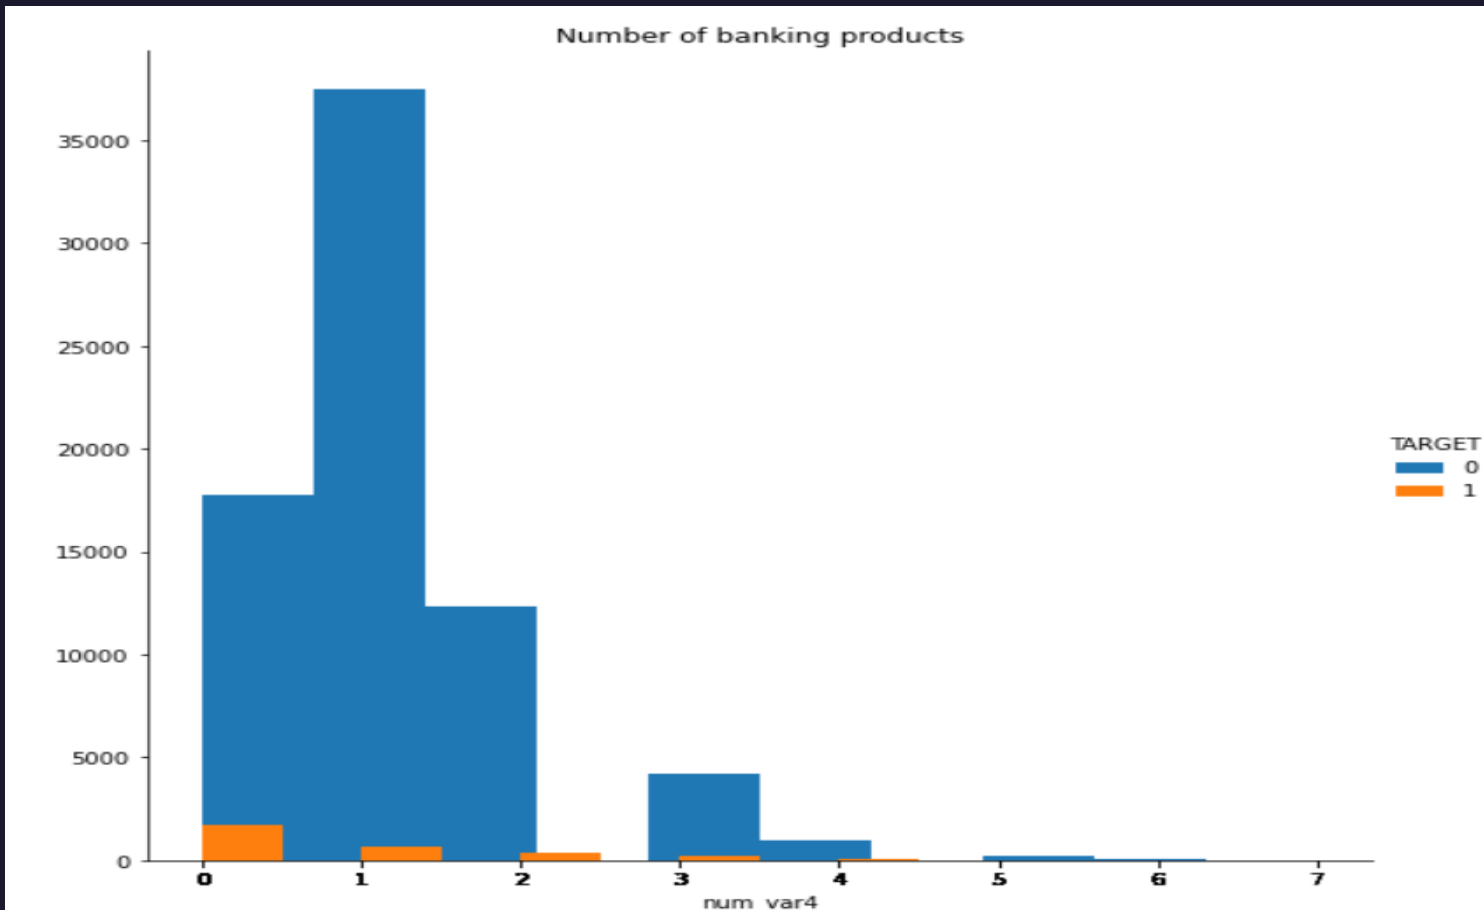


EDA feature 2

```
#EDA
#Feature 2
import seaborn as sns
import matplotlib.pyplot as plt

g=sns.FacetGrid(trainData, hue="TARGET", height=8, aspect=1).map(plt.hist, "num_var4").add_legend()
g.set(xticks=trainData.num_var4)

plt.title("Number of banking products")
plt.show()
```



Field name: num_var4

Range of values : 0 to 9

Probable column: Number of banking products
(‘num_var4’ feature according to some literature is believed to be number of banking products bought by each customer)



EDA feature 2

```
✓ [17] trainData['num_var4'].value_counts()
0s
```

1	38147
0	19528
2	12692
3	4377
4	1031
5	203
6	36
7	6

Name: num_var4, dtype: int64

Field name: num_var4

Range of values : 0 to 9

Probable column: Number of banking products

ANALYSIS :

From the given output, we can analyze that the most preferred banking product is 1 and the least preferred banking products are 6 and 7.

EDA feature 3

```
trainData.loc[trainData['TARGET']==1]['num_var4'].value_counts()
```

0	1737
1	692
2	333
3	182
4	58
5	6

Name: num_var4, dtype: int64

Banking products bought by unsatisfied customers

ANALYSIS :

Number of banking products 6 and 7 didn't make it in the series below.

So, customers who have 6 or more than 6 banking products with the bank are all satisfied.

This could be used to make a new feature while feature engineering whether value of 'num_var4' is 6 or more than 6.



Thank You

Siri Kademani

skademan@asu.edu

